

# 15618 project: GPU object tracking

Denis Merigoux (dmerigou) & Ilai Deutel (ideutel)

Tuesday, April 25<sup>th</sup>

## Abstract

We are going to implement an optimized object tracker on NVIDIA GPUs. The goal is to perform real-time object tracking. Here is the link to our GitHub repository.

## 1 Changes from initial proposal

After having reviewed [2] which presented the SCM algorithm we wanted to implement, we realized that the code provided by the authors wasn't a good starting point (Matlab 2009 code running only on a Windows XP virtual machine). Instead, we chose to take as a starting point the code for the KCF tracker [1] for which a sequential implementation exist in OpenCV. Our new goal is to parallelize this tracker to make it run on the GPU, using OpenCV's CUDA framework. Our work could have more potential and usefulness as part of the OpenCV framework.

## 2 Work completed so far

The first task was to setup a build system, which is not easy given the intricacy of OpenCV, the various options needed to run the tracker or having CUDA support, and the lack of root privileges on the GHC machines. We had to develop various build scripts and try several configurations before reaching an optimal build workflow.

Then we proceeded to implement a correctness test and a profiling analysis of the tracker code, to monitor our progress when parallelizing the code. We test the tracker on a 141-frames fullHD video, and take the average computing time over all frames for the profiling (see figure 1). The analysis of this profiling led us to focus our first efforts into parallelizing the Discrete Fourier transform used by KCF, task still under development since the beginning of the week.

## 3 Updated goals

### 3.1 Primary goals

- Having a correct sequential implementation of the algorithm using C++. **Done.**
- Analyse the workload and determine which bottlenecks we need to parallelize in CUDA. **Done.**
- Parallelize, verify correctness, compute speedup compared to sequential implementation.

| Phase                               | Baseline implemtation time (ms) |
|-------------------------------------|---------------------------------|
| Initialization                      | 308.471                         |
| Average frame time                  | 111.353                         |
| Detection                           | 47.48                           |
| Extract and pre-process the patches | 1.055                           |
| Non-compressed custom descriptors   | 0.193                           |
| Compressed descriptors              | 7.355                           |
| Compressed custom descriptors       | 2.774                           |
| Compress features and KRSL          | 9.372                           |
| Merge all features                  | 1.558                           |
| Compute the gaussian kernel         | 20.087                          |
| Compute the FFT                     | 1.748                           |
| Calculate filter response           | 3.101                           |
| Extract maximum response            | 0.236                           |
| Extracting patches                  | 14.602                          |
| Update bounding box                 | 0                               |
| Non-compressed descriptors          | 1.02                            |
| Non-compressed custom descriptors   | 0.196                           |
| Compressed descriptors              | 7.301                           |
| Compressed custom descriptors       | 3.058                           |
| Update training data                | 3.027                           |
| Feature compression                 | 25.118                          |
| Update projection matrix            | 20.164                          |
| Compress                            | 4.249                           |
| Merge all features                  | 0.705                           |
| Least Squares Regression            | 22.638                          |
| Initialization                      | 0                               |
| Calculate alphas                    | 18.57                           |
| Compute FFT                         | 1.758                           |
| Add a small value                   | 0.378                           |
| New alphaf                          | 1.095                           |
| Update RLS Model                    | 0.837                           |

Figure 1: Breakdown of times for each phase of the KCF tracker’s procedure to update the bouding box for each new frame.

- Analyse the result of our implementation on the Need for speed dataset and compute the equivalent framerate at which our algorithm is capable of tracking the object.

The main grading criteria would be the framerate reached by our parallel algorithm.

### 3.2 Additional goals

- Find or implement an initialization front-end to trace bounding boxes around objects to track.
- Depending on the performance of the algorithm, use our program to actually track objects in real time from a camera feed using the initialization front-end.

### 3.3 Parallelism competition

We may present a graph showing plotting the average time taken to update one frame by the sequential algorithm and our parallelized algorithm, against the resolution of the video (larger resolutions should yield more opportunity for parallelism).

## 4 Revised schedule

| Week          | Task   |
|---------------|--|
| 04/10 – 04/16 | Make a short bibliography about object tracking, build OpenCV and all its dependencies on the GHC machines |
| 04/17 – 04/23 | Analyse performance by timing each step to find the bottlenecks  |
| 04/24 – 04/30 | Implement the parallelized version of the program  |
| 05/01 – 05/06 | Finish optimizing the parallel version and run the benchmarks  |
| 05/07 – 05/12 | If time left, complete the additional goals.   |

## References

- [1] João F. Henriques, Rui Caseiro, Pedro Martins, and Jorge Batista. High-speed tracking with kernelized correlation filters. *CoRR*, abs/1404.7584, 2014.
- [2] W. Zhong, H. Lu, and M. H. Yang. Robust object tracking via sparsity-based collaborative model. In *2012 IEEE Conference on Computer Vision and Pattern Recognition*, pages 1838–1845, June 2012.